

# Nils Matteson

[in](https://linkedin.com/in/nilsmatteson) linkedin.com/in/nilsmatteson [nilsmatteson.com](http://nilsmatteson.com) [nilsmatteson@icloud.com](mailto:nilsmatteson@icloud.com)  Madison, WI

Data Science & CS Senior building **end-to-end ML systems**—from custom streaming infrastructure in Go/Rust to autonomous training pipelines. Experience deploying production RAG systems on AWS. Seeking ML Infrastructure, Backend, or Data Science roles.

## Education

### University of Wisconsin–Madison

B.S. Data Science, Minor in Computer Science

Madison, WI

Expected May 2026

- **Systems & AI:** Big Data Systems (CS 544), Machine Learning (STAT 479), Artificial Intelligence (CS 540), Machine Organization (CS 354), Programming III (CS 400), Intro to Computer Engineering (CS 252).
- **Data Science & Math:** Data Science Modeling I & II (STAT 240/340), DS Programming II (CS 320), Linear Algebra (MATH 340), Discrete Math (MATH 240).

## Technical Skills

**Languages:** Python, Rust, Go, C++, SQL, TypeScript/JavaScript

**ML & Data:** XGBoost, Scikit-learn, PyTorch, Pandas, Feature Engineering, A/B Testing, LLMs, RAG, Hugging Face

**Systems & Cloud:** AWS, GCP, K8s, Docker, gRPC, Kafka, Redis, Distributed Systems, Postgres

**DevOps & Web:** CI/CD, Git, Linux, GitHub Actions, React, Next.js, WebSockets, WebGL

## Experience

### Research Cyberinfrastructure, UW–Madison DoIT

Madison, WI

AI Workflows Research Collaborator

Jan 2026 – Present

- Benchmarked **11 LLMs** on hierarchical RAG pipeline for AI sustainability research; identified scoring bug in boolean outputs, boosting Claude 3.7 Sonnet from 0.570 to **0.721 accuracy** (+26%).
- Built **AWS Bedrock integration** with retry logic, concurrency limits, and throttling handling; optimized indexing from CPU to CUDA, reducing build time from hours to minutes.
- Designed evaluation framework analyzing cost/performance tradeoffs across models; discovered Claude 3.7's edge comes from **citation quality** (0.850 ref overlap vs 0.73 baseline), not raw accuracy.
- Architecting production deployment: Streamlit UI on AWS with S3-hosted SQLite vector indices, comparing against on-prem GB10 GPU cluster for cost/latency analysis.

## Selected Projects

### Madison Metro ML: Autonomous Bus Arrival Prediction

Python, XGBoost, Sentinel, PostgreSQL, React

16K+ LOC end-to-end ML system with ground truth generation, autonomous retraining, and live inference.

- Designed **geospatial ground truth pipeline**: Haversine-based arrival detection matches GPS coordinates to stops (30m threshold), then joins to predictions to compute actual vs. predicted error—solving the “how do we validate transit predictions?” problem.
- Built streaming data pipeline using **Sentinel** (custom message queue) with gRPC producers, dual-threaded consumers with offset tracking, and PostgreSQL persistence for 7K+ daily observations.
- Implemented **autonomous nightly retraining** via GitHub Actions: XGBoost regression on rolling 7-day window with metric-gated deployment (MAE improvement threshold) and Git-versioned model registry.
- Deployed full stack on Railway (collector, consumer, API) and Vercel (React dashboard with MapLibre live tracking); engineered temporal and route-aggregated features with proper train/test split to prevent data leakage.

### Sentinel: Distributed Log Streaming Engine

Go, gRPC, Protobuf, LSM Trees, Raft

Kafka-inspired message queue (5,600+ LOC) powering Madison Metro ML’s real-time data pipeline.

- Engineered custom **LSM-tree storage engine** with skip list memtable achieving 1.7M writes/sec and 3.9M reads/sec; implemented CRC32 checksums, bloom filters, and crash-safe write-ahead log.
- Built **Raft consensus layer** for fault-tolerant leader election and log replication; designed gRPC streaming API with topic/partition semantics, consumer groups, and offset tracking.

### Synapse: Real-time Collaborative Whiteboard

Rust, WASM, WebSockets, CRDTs (Yjs)

Lock-free distributed canvas (Rust + WebAssembly) supporting 50+ concurrent editors.

- Built Rust WebSocket server (Actix) with **CRDT state sync** (Yjs) for conflict-free concurrent editing without operational transforms or central locking.
- Compiled rendering to **WebAssembly** achieving 60 FPS with 10K+ vector objects; architected horizontal scaling via Redis Pub/Sub with session affinity.